

In [12]:

```
%cd /home/mw/project/  
  
/home/mw/project
```

In [13]:

```
import torch  
import torch.nn.functional as F  
import torchvision.transforms as transforms  
import torch.backends.cudnn as cudnn  
import numpy as np  
from nltk.translate.bleu_score import corpus_bleu  
from tqdm import tqdm  
from datasets import *  
from utils import *
```

In [14]:

```
# Parameters  
data_folder = '/home/mw/work/project/coco2014' # folder with data files saved  
data_name = 'coco_5_cap_per_img_5_min_word_freq' # base name shared by data fi  
checkpoint = '/home/mw/project/BEST_checkpoint_coco_5_cap_per_img_5_min_word_fr  
word_map_file = '/home/mw/work/project/coco2014/WORDMAP_coco_5_cap_per_img_5_mi  
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")  
cudnn.benchmark = True  
beam_size = 1  
attention = True
```

In [15]:

```

# Load model
checkpoint = torch.load(checkpoint)
encoder = checkpoint['encoder']
encoder = encoder.to(device)
encoder.eval()
decoder = checkpoint['decoder']
decoder = decoder.to(device)
decoder.eval()

DecoderWithAttention(
  (attention): Attention(
    (encoder_att): Linear(in_features=2048, out_features=512, bias=True)
    (decoder_att): Linear(in_features=512, out_features=512, bias=True)
    (att): Linear(in_features=512, out_features=1, bias=True)
    (relu): ReLU()
    (softmax): Softmax(dim=1)
  )
  (embedding): Embedding(9490, 512)
  (decode_step): LSTMCell(2560, 512)
  (init_h): Linear(in_features=2048, out_features=512, bias=True)
  (init_c): Linear(in_features=2048, out_features=512, bias=True)
  (beta): Linear(in_features=512, out_features=1, bias=True)
  (fc): Linear(in_features=512, out_features=9490, bias=True)
  (dropout_layer): Dropout(p=0.5, inplace=False)
  (bn): BatchNorm1d(512, eps=1e-05, momentum=0.01, affine=True, track_running_st
)

```

In [16]:

```

# Load word map (word2ix)
with open(word_map_file, 'r') as j:
    word_map = json.load(j)
rev_word_map = {v: k for k, v in word_map.items()}
vocab_size = len(word_map)

```

In [17]:

```

# Normalization transform
normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                std=[0.229, 0.224, 0.225])

```

In [18]:

```

# DataLoader
loader = torch.utils.data.DataLoader(
    CaptionDataset(data_folder, data_name, 'TEST', transform=transforms.Compose
    batch_size=1, shuffle=False, num_workers=1, pin_memory=True)
)

```

In [19]:

```
# Lists to store references (true captions), and hypothesis (prediction) for ea  
# If for n images, we have n hypotheses, and references a, b, c... for each ima  
# references = [[ref1a, ref1b, ref1c], [ref2a, ref2b], ...], hypotheses = [hyp1  
references = list()  
hypotheses = list()
```

In [20]:

```
# For each image
for i, (image, caps, caplens, allcaps) in enumerate(tqdm(loader, desc="EVALUATING")):
    k = beam_size

    # Tensor to store top k previous words at each step; now they're just <start>
    k_prev_words = torch.LongTensor([[word_map['<start>']]] * k).to(device) # (k, 1)

    # Tensor to store top k sequences; now they're just <start>
    seqs = k_prev_words # (k, 1)

    # Tensor to store top k sequences' scores; now they're just 0
    top_k_scores = torch.zeros(k, 1).to(device) # (k, 1)

    # Lists to store completed sequences and scores
    complete_seqs = list()
    complete_seqs_scores = list()

    # Move to GPU device, if available
    image = image.to(device) # (1, 3, 256, 256)

    # Encode
    encoder_out = encoder(image) # (1, enc_image_size, enc_image_size, encoder_dim)

    # Flatten encoding
    # We'll treat the problem as having a batch size of k
    if attention:
        encoder_dim = encoder_out.size(3)
        encoder_out = encoder_out.view(1, -1, encoder_dim) # (1, num_pixels, encoder_dim)
        num_pixels = encoder_out.size(1)
        encoder_out = encoder_out.expand(k, num_pixels, encoder_dim) # (k, num_pixels, encoder_dim)
    else:
        encoder_out = encoder_out.reshape(1, -1)
        encoder_dim = encoder_out.size(1)
        encoder_out = encoder_out.expand(k, encoder_dim)

    # Start decoding
    step = 1
    if attention:
        mean_encoder_out = encoder_out.mean(dim=1)
        h = decoder.init_h(mean_encoder_out) # (1, decoder_dim)
        c = decoder.init_c(mean_encoder_out)
    else:
        init_input = decoder.bn(decoder.init(encoder_out))
        h, c = decoder.decode_step(init_input) # (batch_size_t, decoder_dim)

    smoth_wrong = False

    # s is a number less than or equal to k, because sequences are removed from
```

```

while True:
    embeddings = decoder.embedding(k_prev_words).squeeze(1) # (s, embed_dim)
    if attention:
        scores, _, h, c = decoder.one_step(embeddings, encoder_out, h, c)
    else:
        scores, h, c = decoder.one_step(embeddings, h, c)
    scores = F.log_softmax(scores, dim=1)
    scores = top_k_scores.expand_as(scores) + scores # (s, vocab_size)
    # For the first step, all k points will have the same scores (since same)
    if step == 1:
        top_k_scores, top_k_words = scores[0].topk(k, 0, True, True) # (s)
    else:
        # Unroll and find top scores, and their unrolled indices
        top_k_scores, top_k_words = scores.view(-1).topk(k, 0, True, True)
        # Convert unrolled indices to actual indices of scores
        prev_word_inds = top_k_words // vocab_size # (s)
        next_word_inds = top_k_words % vocab_size # (s)
        # Add new words to sequences
        seqs = torch.cat([seqs[prev_word_inds], next_word_inds.unsqueeze(1)], dim=1)
        # Which sequences are incomplete (didn't reach <end>)?
        incomplete_inds = [ind for ind, next_word in enumerate(next_word_inds)
                           if next_word != word_map['<end>']]
        complete_inds = list(set(range(len(next_word_inds)) - len(incomplete_inds)))
        # Set aside complete sequences
        if len(complete_inds) > 0:
            complete_seqs.extend(seqs[complete_inds].tolist())
            complete_seqs_scores.extend(top_k_scores[complete_inds])
        k -= len(complete_inds) # reduce beam length accordingly
        # Proceed with incomplete sequences
        if k == 0:
            break
        seqs = seqs[incomplete_inds]
        h = h[prev_word_inds[incomplete_inds]]
        c = c[prev_word_inds[incomplete_inds]]
        encoder_out = encoder_out[prev_word_inds[incomplete_inds]]
        top_k_scores = top_k_scores[incomplete_inds].unsqueeze(1)
        k_prev_words = next_word_inds[incomplete_inds].unsqueeze(1)
        # Break if things have been going on too long
        if step > 50:
            smoth_wrong = True
            break
        step += 1
    if not smoth_wrong:
        i = complete_seqs_scores.index(max(complete_seqs_scores))
        seq = complete_seqs[i]
    else:
        seq = seqs[0][:20]
    # References
    img_caps = allcaps[0].tolist()

```

```

img_captions = list(
    map(lambda c: [w for w in c if w not in {word_map['<start>'], word_map[
        img_caps)}] # remove <start> and pads
references.append(img_captions)
# Hypotheses
hypotheses.append([w for w in seq if w not in {word_map['<start>'], word_ma
assert len(references) == len(hypotheses)

```

```

EVALUATING AT BEAM SIZE 1: 0%|          | 0/25000 [00:00<?, ?it/s]/opt/conda/l
EVALUATING AT BEAM SIZE 1: 100%|██████████| 25000/25000 [14:08<00:00, 29.46it/s]

```

In [21]:

```

# Calculate BLEU-4 scores
bleu4 = corpus_bleu(references, hypotheses)
print(bleu4)

```

```
0.2607194479996266
```

In []: